

Inference Benchmark Report

Generated: Thu Sep 25 04:37:49 2025

Overview

This report summarizes inference accuracy, latency and throughput across different backends (PyTorch CPU, PyTorch GPU, TensorRT FP32/FP16/INT8 where available).

Results Summary

pytorch_cpu

- Accuracy: **0.0560**
- Latency per image (ms): 14.45
- Throughput (img/s): 1081.40

pytorch_gpu

- Accuracy: **0.0560**
- Latency per image (ms): 0.11
- Throughput (img/s): 141158.22

trt_fp32

- Error: `tensorrt\_bindings.tensorrt.Builder` object has no attribute 'build_engine`

trt_fp16

- Error: `tensorrt\_bindings.tensorrt.Builder` object has no attribute 'build_engine`

trt_int8

- Error: `tensorrt\_bindings.tensorrt.Builder` object has no attribute 'build_engine`

Automated Observations

- GPU vs CPU speedup (latency): ~130.53x

Plots and Confusion Matrices

- `cm\_cpu.png`, `cm\_gpu.png`, `cm\_trt\_fp32.png`, etc. contain confusion matrices.
- `mis\_cpu.png`, `mis\_gpu.png`, `mis\_trt\_fp32.png`, etc. contain example misclassified images.









